Query Latency Comparison: ZGQ vs HNSW

Legend: 10K vectors, 100K vectors

Trade-off: 3.1x slower for 20% memory savings

HNSW: 0.0152 (10K vectors), 0.0453 (100K vectors)

ZGQ (Proposed): 0.0620 (10K vectors), 0.1397 (100K vectors)

X-axis: Algorithm

Y-axis: Query Latency (ms)